

Setting YARN NodeManager Labels in Airflow DAGs

- Triggering spark job with nodelabels option ensures your job to land on particular YARN nodemanager
- Set nodelabels only if your DAG consists atleast one or more `SparkSubmitOperator`
- For each new DAG creation , notify Infra team to create new YARN nodemanager VM and to map nodelabels to YARN nodemanager Service
- Nodelabel would be same as DAG name but with suffix `nodelabels_`.

For Example: If DAG name is `spark_pi_dag` then `nodelabel` would be `nodelabel_spark_pi_dag`

```

13 Infra_nodeLabel_test = DAG(
14     "spark_pi_dag",
15     description="YARN NodeLabel Testing",
16     schedule_interval="55 4 * * *",
17     start_date=start_timestamp,
18     tags=['Infra', 'infra'],
19     # Internally handled by state management table
20     catchup=False
21 )
22
23 dummy_start = DummyOperator(
24     task_id='start',
25     dag=Infra_nodeLabel_test
26 )
27
28 dummy_end = DummyOperator(
29     task_id='end',
30     dag=Infra_nodeLabel_test
31 )
32
33 start_nodemangers= BashOperator(
34     task_id="start_nodemanager",
35     bash_command="source /opt/python3_6-virtual_envs/airflow_env/bin/activate && cd /data/1/infra/jobs/infra_processing/saas-stage1-monitoring/node-manager-sta
36 )
37
38 stop_nodemangers= BashOperator(
39     task_id="stop_nodemanager",
40     bash_command=" source /opt/python3_6-virtual_envs/airflow_env/bin/activate && cd /data/1/infra/jobs/infra_processing/saas-stage1-monitoring/node-manager-st
41     trigger_rule=TriggerRule.ALL_DONE
42 )
43
44 pi_value_calculation_1 = SparkSubmitOperator(
45     name="airflow_spark_pi_1",
46     application="/opt/spark/examples/jars/spark-examples_2.11-2.3.1.jar",
47     task_id="pi_value_calculation_1",
48     java_class="org.apache.spark.examples.SparkPi",
49     conn_id='spark_default',
50     driver_memory="512M",
51     num_executors=1,
52     executor_cores=1,
53     executor_memory="512M",
54     conf={"spark.yarn.am.nodeLabelExpression": "nodelabel_spark_pi_dag" "spark.yarn.executor.nodeLabelExpression": "nodelabel_spark_pi_dag"}
55 )

```

- Notify Infra team to create a config-file where YARN nodemanager VM details will be placed.

The config file to be placed in DAG folder path.

```
(airflow_env) [airflow@argoid-saas-stage-host-012 ~]$
(airflow_env) [airflow@argoid-saas-stage-host-012 ~]$ ls -l ~/airflow/dags/infra/
total 16
-rwxrwxrwx. 1 argoid argoid 961 Nov 3 05:27 airflow-archived-logs.py
-rwxrwxrwx. 1 argoid argoid 1185 Nov 2 16:24 airflow-backup-test.py
drwxrwxrwx. 2 argoid argoid 202 Dec 21 12:37 airflow-backup-test
-rw-r--r--. 1 airflow airflow 156 Dec 21 12:35 yarn-label-test-config.json
-rw-rw-r--. 1 airflow airflow 3189 Dec 21 12:24 yarn-label-test-dag.py
(airflow_env) [airflow@argoid-saas-stage-host-012 ~]$
(airflow_env) [airflow@argoid-saas-stage-host-012 ~]$
(airflow_env) [airflow@argoid-saas-stage-host-012 ~]$
(airflow_env) [airflow@argoid-saas-stage-host-012 ~]$ cat ~/airflow/dags/infra/yarn-label-test-config.json
{
  "nodemanager_host_list": ["argoid-saas-stage-host-041"],
  "azure_resource_group": "ARGOID-RESOURCE-GROUP-7-SAAS-STAGE1",
  "health_check_port_url": "8042"
}
(airflow_env) [airflow@argoid-saas-stage-host-012 ~]$
(airflow_env) [airflow@argoid-saas-stage-host-012 ~]$
```

- Add `BashOperator` tasks in your DAG to start and stop Nodemanager VMs also, `trigger_rule` also to be set in `stop_nodemanagers` part.

```

dummy_end = DummyOperator(
    task_id='end',
    dag=Infra_nodeLabel_test)

start_nodemangers= BashOperator(
    task_id= 'start_nodemanager',
    bash_command='source /opt/python3.6-virtual_envs/airflow_env/bin/activate && cd /data/1/infra/jobs/infra_processing/saas-stage1-monitoring/node-manager-start-stop-automation/ && source azure_credentials.sh & python start_nodemanager.py /home/airflow/airflow/dags/infra/yarn-label-test-config.json '
)

stop_nodemangers= BashOperator(
    task_id= 'stop_nodemanager',
    bash_command='source /opt/python3.6-virtual_envs/airflow_env/bin/activate && cd /data/1/infra/jobs/infra_processing/saas-stage1-monitoring/node-manager-start-stop-automation/ && source azure_credentials.sh & python stop_nodemanager.py /home/airflow/airflow/dags/infra/yarn-label-test-config.json ' ,
    trigger_rule=TriggerRule.ALL_DONE
)

pi_value_calculation_1 = SparkSubmitOperator(
    name='airflow_spark_pi_1',
    application='/opt/spark/examples/jars/spark-examples_2.11-2.3.1.jar',
    task_id='pi_value_calculation_1',
    java_class='org.apache.spark.examples.SparkPi',
    conn_id='spark_default',
    driver_memory='512M',
    num_executors=1,
    executor_cores=1,
    executor_memory='512M',
    conf={'spark.yarn.am.nodeLabelExpression': 'nodelabel_spark_pi_dag', 'spark.yarn.executor.nodeLabelExpression': 'nodelabel_spark_pi_dag'})
)

```

- Set nodelabels for `spark.yarn.am.nodeLabelExpression` `spark.yarn.executor.nodeLabelExpression` properties in each `SparkSubmitOperator` tasks

```

pi_value_calculation_2 = SparkSubmitOperator(
    name='airflow_spark_pi_2',
    application='/opt/spark/examples/jars/spark-examples_2.11-2.3.1.jar',
    task_id='pi_value_calculation_2',
    java_class='org.apache.spark.examples.SparkPi',
    conn_id='spark_default',
    driver_memory='512M',
    num_executors=1,
    executor_cores=1,
    executor_memory='512M',
    conf={'spark.yarn.am.nodeLabelExpression': 'nodelabel_spark_pi_dag', 'spark.yarn.executor.nodeLabelExpression': 'nodelabel_spark_pi_dag'})

pi_value_calculation_3 = SparkSubmitOperator(
    name='airflow_spark_pi_3',
    application='/opt/spark/examples/jars/spark-examples_2.11-2.3.1.jar',
    task_id='pi_value_calculation_3',
    java_class='org.apache.spark.examples.SparkPi',
    conn_id='spark_default',
    driver_memory='512M',
    num_executors=1,
    executor_cores=1,
    executor_memory='512M',
    conf={'spark.yarn.am.nodeLabelExpression': 'nodelabel_spark_pi_dag', 'spark.yarn.executor.nodeLabelExpression': 'nodelabel_spark_pi_dag'})

dummy_start >> start_nodemangers >> pi_value_calculation_1 >> pi_value_calculation_2 >> pi_value_calculation_3 >> stop_nodemangers >> dummy_end

```

- Keep `stop_nodemangers` task to run just before the `dummy_end` task

```

    driver_memory='512M',
    num_executors=1,
    executor_cores=1,
    executor_memory='512M',
    conf={'spark.yarn.am.nodeLabelExpression': 'nodelabel_spark_pi_dag', 'spark.yarn.executor.nodeLabelExpression': 'nodelabel_spark_pi_dag'})
)

dummy_start >> start_nodemangers >> pi_value_calculation_1 >> pi_value_calculation_2 >> pi_value_calculation_3 >> stop_nodemangers >> dummy_end

```